

# Django Production-Readiness Assignment

Blue-green deploys • WebSockets • Observability

## 1. Context & Goal

We are rolling a **brand-new WebSocket service** into a Django-based stack.

Your task is to deliver a **self-contained repo** that:

- implements the WebSocket feature itself
- demonstrates you understand the ASGI concurrency model (event loop vs. thread-pool work)
- enables zero-downtime blue/green releases with Docker
- provides runtime observability: structured logs, metrics, health checks & alerts

The feature logic is intentionally simple (a message counter over WebSockets) so you can focus on production concerns.

## 2. Functional Requirements

Area	Requirement
<b>WebSocket endpoint</b>	<code>/ws/chat/</code> accepts a text message; replies with <code>{"count": n}</code> where <code>n</code> is the number of messages seen on that connection. On close, send <code>{"bye": true, "total": n}</code> and clean up.
<b>Server-side push</b>	Every 30 s broadcast a heartbeat <code>{"ts": &lt;iso-timestamp&gt;}</code> to all active sockets.

Area	Requirement
<b>Graceful shutdown</b>	When the container receives <code>SIGTERM</code> , finish in-flight messages, close sockets with code 1001, and exit within 10 s.
<b>Reconnection (nice-to-have)</b>	A reconnecting client may include its previous session UUID as a query param; resume the counter if still in memory.

### 3. Production-grade Requirements

Theme	Must-have	Nice-to-have / Suggestions
<b>Blue-green</b>	Two independent colour stacks ( <code>app_blue</code> , <code>app_green</code> ) behind a reverse proxy. A promotion script (bash or Make target) should: 1) build & start the next colour, 2) run smoke tests, 3) flip traffic, 4) retire the old colour.	Use Docker Compose with Nginx or Traefik as the router. If you prefer k8s, a single-node Kind cluster is fine.
<b>Concurrency tuning</b>	Show how many ASGI workers and thread-pool workers you run, and why (CPU-bound vs I/O). Provide a comment in <code>docker-compose.yml</code> or <code>helm/values.yml</code> .	Use <code>uvicorn --workers</code> , <code>--loop</code> , and Django Channels' <code>ASYNC_CAPABLE</code> docs.

Theme	Must-have	Nice-to-have / Suggestions
<b>Observability</b>	<b>Metrics</b> endpoint <code>/metrics</code> exposing Prometheus-style counters: total messages, active connections, error count, shutdown time.	

**Structured logs** (JSON) for each event with `request_id`.

**Health probes:** `/healthz` (liveness) + `/readyz` (readiness toggled on startup/shutdown). | Wire Grafana dashboards via a pre-built `grafana/dashboards` JSON; or ship logs to Loki. |

**Alerting rule** | Create one Prometheus rule (YAML) that triggers if active connections drop to 0 for >60 s. | **Monitoring script** | A bash/py script `monitor.sh` that: 1) tails container logs for `ERROR`, 2) hits `/metrics` and prints the top-5 counters every 10 s. | **CI task** | GitHub Actions or GitLab CI pipeline that: build-&-test, run `pytest`, execute the monitor script for 20 s, and archive logs/artifacts. |

## 4. Non-Functional Targets

- **Throughput:** Demonstrate  $\geq 5\,000$  concurrent sockets on your laptop.
- **Startup** < 3 s; **shutdown** < 10 s.

## 5. Deliverables

```
.
├─ docker/
│   └─ Dockerfile
│       └─ compose.yml
├─ app/
│   └─ manage.py
```

```
|   ├── settings.py
|   └── chat/
├── scripts/
|   ├── promote.sh
|   └── monitor.sh
├── ci/
└── docs/
    ├── DESIGN.md
    └── OBSERVABILITY.md
```

**README** must include:

- the exact one-liner to spin up the full stack
- how to run the simple Locust/Autobahn load test ( `make load-test` )
- how to flip blue→green

Optionally, add a short 5-minute Loom or GIF walking through the deploy & metrics.

## 6. Suggested Building Blocks (optional)

Concern	Python / Django ecosystem pick
ASGI server	<code>uvicorn[standard]</code> or <code>daphne</code>
WebSockets	<code>django-channels</code> 4.x
Thread-pool offload	<code>asgiref.sync.sync_to_async</code> & <code>ThreadPoolExecutor</code>
Metrics	<code>prometheus-client</code> or <code>django-prometheus</code>
Structured logs	<code>python-json-logger</code> or <code>structlog</code>
Graceful shutdown	<code>lifespan</code> hooks in Uvicorn, or Django's <code>request_finished</code> signal
Smoke tests	<code>pytest-asyncio</code> , <code>websockets</code> library

Concern	Python / Django ecosystem pick
Load test	Locust , k6 , or a minimal asyncio script
Reverse proxy	nginx:alpine , traefik:v3

Feel free to swap for equivalents you know well; just document the choice.

## 7. Evaluation Checklist

Weight	Criterion	Details
40 %	<b>Production architecture</b>	Correct blue/green flow, health probes, graceful shutdown, Docker layering, config separation.
25 %	<b>ASGI &amp; concurrency insight</b>	Proper use of workers vs thread-pool, no shared mutable state pitfalls, clear explanation in DESIGN.md.
15 %	<b>Observability</b>	Metrics coverage, log quality, alert rule validity, demo dashboards.
10 %	<b>Code quality &amp; tests</b>	Idiomatic Django, typing where useful, pytest coverage.
10 %	<b>Docs &amp; Dev-X</b>	make dev-up runs in $\leq 3$ commands, README clarity, comments on trade-offs.

Bonus: reconnection support.

## 8. Submission

Create a repo, push your code, and email **Yugesh** ( `yugesh.kothari@sigiq.ai` ) or WhatsApp **+41 762621416** when ready. We'll schedule a review call to walk through your solution and discuss technical decisions.